

# **BCSE498J Project-II– Capstone Project**

## **CONVERTING BUSINESS QUERIES TO VALUABLE INSIGHTS AND VISUALIZATIONS USING GEN AI**

*Submitted in partial fulfillment of the requirements for the degree of*

**Bachelor of Technology**

*in*

**Computer Science and Engineering**

*by*

**21BCE0525    ROHITH PARAHMESH K**

**21BCE2624    GOLI ANSHUL REDDY**

**Under the Supervision of**

**Dr. Santhi H**

Associate Professor Sr

School of Computer Science and Engineering (SCOPE)



April 2025

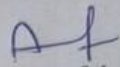
### DECLARATION

I hereby declare that the project entitles "Converting Business Queries to Valuable Insights and Visualizations Using Gen AI" submitted by me, for the award of the degree of *Bachelor of Technology in Computer Science and Engineering* to VIT is a record of bonafide work carried out by me under the supervision of Dr. Santhi H.

I further declare that the work reported in this project has not been submitted and will not be submitted, either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university.

Place : Vellore

Date : 23-04-2025

  
Signature of the Candidate

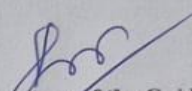
### CERTIFICATE

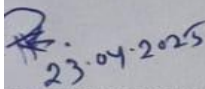
This is to certify that the project entitled "Converting Business Queries to Valuable Insights and Visualizations Using Gen AI" submitted by Rohith Parahmesh K (21BCE0525) and Goli Anshul Reddy (21BCE2624), School of Computer Science and Engineering, VIT , for the award of the degree of *Bachelor of Technology in Computer Science and Engineering*, is a record of bonafide work carried out by him / her under my supervision during Winter Semester 2024-2025, as per the VIT code of academic and research ethics.

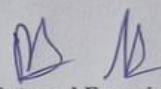
The contents of this report have not been submitted and will not be submitted either in part or in full, for the award of any other degree or diploma in this institute or any other institute or university. The project fulfills the requirements and regulations of VIT University and in my opinion meets the necessary standards for submission.

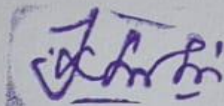
Place : Vellore

Date : 23-04-25

  
Signature of the Guide

  
23.04.2025  
Internal Examiner

  
External Examiner



Dr. UMADEVI K S

Computer Science and Engineering

## EXECUTIVE SUMMARY

This project introduces a framework that leverages Generative AI to convert business-specific natural language queries into insightful analytics and visualizations. In an increasingly data-driven landscape, businesses are often hindered by the complexity of data querying and analysis tools. Our approach bridges this gap using AI-powered interfaces that allow non-technical users to gain valuable insights by simply asking questions in natural language.

The system combines Natural Language Processing (NLP), Machine Learning (ML), and visualization techniques to interpret user queries, extract relevant datasets, conduct predictive analysis, and generate dynamic dashboards. It supports adaptive querying, handles ambiguous phrasing, and ensures accurate visual output using context-aware processing. The platform's strength lies in its scalability, domain adaptability, and explainability.

This solution democratizes access to business intelligence by reducing dependency on data analysts and lowering the technical barriers for decision-making. It can be used across industries such as finance, marketing, operations, and customer engagement. The system is designed with modular components including a query processor, context resolver, predictor, and visualization engine. Each component contributes to seamless transformation from business questions to meaningful output.

Overall, this project showcases a transformative method of engaging with data through conversational AI, enhancing both efficiency and decision quality in real-time. It lays the groundwork for future expansion into real-time data handling, enhanced visualization interactivity, and domain-specific NLP fine-tuning.

## ACKNOWLEDGEMENTS

I am deeply grateful to the management of Vellore Institute of Technology (VIT) for providing me with the opportunity and resources to undertake this project. Their commitment to fostering a conducive learning environment has been instrumental in my academic journey. The support and infrastructure provided by VIT have enabled me to explore and develop my ideas to their fullest potential.

My sincere thanks to Dr. Jaisankar N, Dean - School of Computer Science and Engineering (SCOPE), for his unwavering support and encouragement. His leadership and vision have greatly inspired me to strive for excellence.

I express my profound appreciation Dr. Umadevi K, the Head of the Software Systems, for his/her insightful guidance and continuous support. His/her expertise and advice have been crucial in shaping throughout the course. His/her constructive feedback and encouragement have been invaluable in overcoming challenges and achieving goals.

I am immensely thankful to my project supervisor, Dr.Santhi H, for his dedicated mentorship and invaluable feedback. His patience, knowledge, and encouragement have been pivotal in the successful completion of this project. My supervisor's willingness to share his expertise and provide thoughtful guidance has been instrumental in refining my ideas and methodologies. His support has not only contributed to the success of this project but has also enriched my overall academic experience.

Thank you all for your contributions and support.

Rohith Parahmesh K  
Goli Anshul Reddy

— Af

## TABLE OF CONTENTS

Sl.No	Contents	Page No.
	<b>Acknowledgement</b>	<b>iii</b>
	<b>Executive Summary</b>	<b>iv</b>
	<b>List of Tables</b>	<b>vii</b>
	<b>List of Figures</b>	<b>viii</b>
	<b>Abbreviations</b>	<b>ix</b>
	<b>Symbols and Notations</b>	<b>x</b>
<b>1.</b>	<b>INTRODUCTION</b>	<b>1</b>
	1.1 BACKGROUND	1
	1.2 MOTIVATIONS	2
	1.3 SCOPE OF THE PROJECT	3
<b>2.</b>	<b>PROJECT DESCRIPTION AND GOALS</b>	<b>5</b>
	2.1 LITERATURE REVIEW	5
	2.2 RESEARCH GAP	8
	2.3 OBJECTIVES	9
	2.4 PROBLEM STATEMENT	10
	2.5 PROJECT PLAN	11
<b>3.</b>	<b>TECHNICAL SPECIFICATION</b>	<b>14</b>
	3.1 REQUIREMENTS	14
	3.1.1 Functional	14
	3.1.2 Non-Functional	15
	3.2 FEASIBILITY STUDY	16
	3.2.1 Technical Feasibility	16
	3.2.2 Economic Feasibility	17
	3.2.2 Social Feasibility	17
	3.3 SYSTEM SPECIFICATION	18
	3.3.1 Hardware Specification	18
	3.3.2 Software Specification	19
<b>4.</b>	<b>DESIGN APPROACH AND DETAILS</b>	<b>21</b>
	4.1 SYSTEM ARCHITECTURE	21
	4.2 DESIGN	23
	4.2.1 Data Flow Diagram	23
	4.2.2 Sequence Diagram	25

<b>5.</b>	<b>METHODOLOGY AND TESTING</b>	<b>26</b>
	5.1 Module Description	26
	5.2 Testing	29
<b>6.</b>	<b>PROJECT DEMONSTRATION</b>	<b>31</b>
<b>7.</b>	<b>RESULT AND DISCUSSION</b>	<b>33</b>
	7.1 System performance and functional Achievements	33
	7.2 User Experience and Usability	34
	7.3 Comparison with Existing Research and Tools	35
	7.4 Innovation and Novel Contributions	36
	7.5 Challenges and Limitations	37
	7.6 Summary of Impact	
<b>8.</b>	<b>CONCLUSION AND FUTURE ENHANCEMENTS</b>	<b>38</b>
	8.1 Conclusion	38
	8.2 Future Enhancements	38
<b>9.</b>	<b>INDIVIDUAL CONTRIBUTION</b>	<b>40</b>
<b>10.</b>	<b>REFERENCES</b>	<b>41</b>
	<b>APPENDIX A – SAMPLE CODE</b>	<b>43</b>

## **List of Tables**

<b>Table No.</b>	<b>Title</b>	<b>Page No.</b>
1.	Literature Review/Gaps identified	5
2	Project Outcome and differences	44



## **List of Figures**

<b>Figure No.</b>	<b>Title</b>	<b>Page No.</b>
1.	Working Project Plan	8
2.	Proposed System Architecture	21
3.	Data Flow Diagram Level 0	22
4.	Sequence Diagram of Working Model	23
5.	ScreenShot of Model UI	31
6.	ScreenShot of Insight Generation	32

## **List of Abbreviations**

AI - Artificial Intelligence

BI - Business Intelligence

CSV - Comma-Separated Values

LLM - Large Language Model

NLP - Natural Language Processing

NL2VIS - Natural Language to Visualization

IQR - Interquartile Range

SQL - Structured Query Language

## Symbols and Notations

$\mu$  - Mean (average value)

$\sigma$  - Standard deviation

IQR - Interquartile Range

Q1 - First Quartile

Q3 - Third Quartile

$\rightarrow$  - Indicates process flow or transformation

$\Sigma$  - Summation

## Chapter 1

# INTRODUCTION

### 1.1 BACKGROUND

Generative AI and Natural Language Processing (NLP) have rapidly evolved, influencing fields such as business intelligence, education, and finance. With the increasing reliance on data-driven decision-making, organizations require tools that simplify complex data interactions, allowing non-technical users to extract meaningful insights without the need for advanced SQL or programming skills. The emergence of neural databases and AI-powered data retrieval systems has addressed some of these challenges, yet significant limitations persist, particularly in understanding ambiguous queries, handling domain-specific contexts, and providing explainable results (\*\*Thorne et al., 2021\*\*).

Traditional business intelligence tools often rely on structured queries and predefined templates, limiting flexibility in data exploration. The shift toward conversational AI-driven interfaces aims to bridge this gap by allowing users to interact with databases using natural language. However, current implementations struggle with contextual understanding, ambiguity resolution, and real-time adaptability. These challenges highlight the need for more robust AI-driven query processing frameworks that can interpret, analyze, and visualize data dynamically (\*\*Bandi et al., 2023\*\*).

Additionally, ethical concerns surrounding AI-generated insights, particularly regarding bias, reliability, and explainability, remain a critical area of focus (\*\*Zohny et al., 2023\*\*). While AI has demonstrated significant advancements in content generation and predictive analytics, its effectiveness in business intelligence applications depends on transparency, interpretability, and accuracy. Addressing these concerns requires a holistic approach that integrates generative AI capabilities with trustworthy, explainable, and domain-adaptive mechanisms.

### 1.2 MOTIVATIONS

Despite the significant advancements in AI-driven business intelligence, several research gaps hinder its widespread adoption. Existing studies have explored Generative AI applications in finance and decision support (\*\*Chen et al., 2023\*\*) but have noted challenges in real-time data processing, adaptability, and accuracy. Similarly, \*\*Johri et al. (2023)\*\* identified AI's potential in education but emphasized its limitations in contextual reasoning and user engagement. These findings underscore the necessity of developing AI-driven interfaces that enhance human-AI collaboration by improving query interpretation, personalization, and domain specificity.

Furthermore, while research has extensively covered evaluation metrics, model

architectures, and input-output formats of Generative AI (\*\*Bandi et al., 2023\*\*), the practical implementation of explainable, business-focused AI tools remains limited. Current systems often fail to contextualize queries within dynamic business environments, leading to inaccurate or non-actionable insights. This lack of adaptability reduces trust in AI-driven decision-making, further limiting its potential impact on business operations.

The ethical dimensions of Generative AI also play a crucial role in its acceptance. Issues such as misuse, bias, and lack of transparency in AI-generated content pose significant risks (\*\*Megahed et al., 2023\*\*). In response to these challenges, our research aims to bridge these gaps by integrating machine learning-driven predictive analytics with an interactive, explainable, and user-friendly AI interface. By enhancing query adaptability, visualization capabilities, and interpretability, we strive to make business intelligence more accessible, efficient, and trustworthy.

### **1.3 SCOPE OF THE PROJECT**

The scope of this project is to design and develop an AI-driven framework that enables users to transform business-related natural language queries into actionable insights and visualizations, without requiring any prior technical or programming knowledge. In traditional analytics workflows, significant expertise is needed to interact with databases, construct queries, interpret results, and generate visuals. This project aims to bridge that gap by leveraging Generative AI, Natural Language Processing (NLP), and data visualization techniques.

The framework accommodates various business domains such as finance, operations, marketing, and logistics, making it a versatile and domain-agnostic tool. It processes user queries in plain English, analyzes the context, extracts relevant data, and generates visual dashboards or summaries. The backend architecture includes a language model-driven parser, a data fetcher, a prediction module for trend analysis, and a visualization engine. These components interact seamlessly to provide a real-time response to the user's input.

This solution democratizes data access and supports better, faster decision-making by enabling stakeholders—including those without technical backgrounds—to independently explore business data. The system's flexibility allows it to scale across

departments and integrate with existing business intelligence pipelines or APIs.

The project also includes considerations for ethical AI usage, focusing on interpretability, transparency, and bias mitigation in the generation of insights. By providing a clear explanation of how a query was processed and how insights were derived, the platform fosters trust among its users.

From a technical standpoint, the scope involves building a front-end interface for query input and result visualization, integrating it with a backend powered by pre-trained language models (e.g., GPT), and deploying the system on a local or cloud environment for testing. Functional testing, user feedback collection, and performance optimization are integral parts of the development lifecycle.

In essence, this project transforms the way business users interact with data—moving from static, manual processes to dynamic, conversational experiences powered by AI. It sets the foundation for future enhancements including real-time analytics, voice-based queries, multilingual support, and seamless integration with enterprise tools such as CRM and ERP systems.

## Chapter 2

# PROJECT DESCRIPTION AND GOALS

## 2.1 LITERATURE REVIEW

*Table 1: Literature Review*

Study	Contribution	Methodology	Gaps Identified
Yu et al.	NLP-based	Translates user queries	Limited support for

(FlowSense, 2023)	dataflow visualization system	into visual representations	complex queries, lacks explainability, does not leverage recent LLM advancements
Li Wang et al. (2023)	Evaluates LLM-based visualization generation	Assesses accuracy of NL2VIS translation and error analysis	Misinterpretation of data attributes, syntax errors in generated specifications
Gao & Dontcheva (DataTone, 2022)	Manages ambiguity in NL-based data visualization	Interactive clarification and multiple interpretations	Scalability and generalization challenges, lacks integration with advanced LLMs
Survey: Towards Natural Language Interfaces for Data Visualization	Reviews Visualization-oriented Natural Language Interfaces (V-NLIs)	Categorizes NLIs by query interpretation, data transformation, and visual mapping	Lacks discussion on the impact of recent LLMs on V-NLIs
Survey: Natural Language Interfaces for Tabular Data Querying and Visualization	Examines NL interfaces for querying and visualizing tabular data	Focuses on semantic parsing and LLM impact	Needs more exploration of real-world deployment challenges
Mirror (2023)	NL interface for database querying, summarization, and visualization	Generates SQL queries from natural language and provides visual feedback	Struggles with handling highly ambiguous or multi-step queries
Talk2Data (2021)	Exploratory data analysis via NL question decomposition	Breaks complex queries into sub-questions for step-by-step analysis	Limited scalability and struggles with diverse analytical question types

## 2.2 RESEARCH GAP

Despite significant advancements in AI-driven NL2VIS systems, key challenges persist in real-world applications. Many models struggle with handling complex analytical

queries, particularly those involving multi-step reasoning and ambiguous user intents. Explainability remains a major gap, as most frameworks fail to clearly articulate how user queries are interpreted and transformed into visualizations.

Additionally, misinterpretation of data attributes and syntax errors in automatically generated visualization specifications continue to hinder usability. Many systems lack scalability and real-time adaptability, making them impractical for dynamic business intelligence and large-scale data analysis.

Future research should focus on enhancing real-time query interpretation, improving ambiguity resolution, and integrating advanced Large Language Models (LLMs) to refine natural language understanding. Expanding support for multi-variable relationships and interactive visualization generation will further improve the practicality of AI-powered NL2VIS solutions.

## **2.3 OBJECTIVES**

The core objective of this project is to bridge the gap between non-technical business stakeholders and the increasingly complex world of data analytics. In the current digital era, businesses generate enormous amounts of data across various departments—finance, marketing, operations, logistics, and customer support. However, making sense of this data often requires technical skills like SQL querying, scripting, and the ability to interpret charts. This project aims to eliminate that barrier by enabling users to interact with data in natural language and obtain actionable insights automatically through AI.

The project aspires to build a robust system that integrates Natural Language Processing (NLP), Machine Learning (ML), and data visualization to offer a seamless analytics experience. One of the main objectives is to develop a natural language interface that accurately understands and processes business-related queries, including those that are ambiguous or loosely structured. This means the system must be capable of recognizing intent, context, and relationships between entities within the query.

Another objective is to establish a backend processing pipeline that can map the interpreted query to the corresponding database or dataset, extract relevant information, and apply necessary operations such as filtering, sorting, or forecasting using ML models. This pipeline must ensure consistency, accuracy, and speed in data retrieval and analysis.

A third major goal is to create a visualization layer that can present results in the form of easy-to-understand and interactive visual formats such as bar charts, pie charts, line graphs, heatmaps, and dashboards. These visualizations should adapt dynamically to user queries and allow further interaction or drill-down capabilities.

Additional objectives include ensuring system transparency and interpretability. Users should be able to understand how their queries were interpreted and how the insights were derived. This fosters trust and ensures the responsible use of AI.

From a technical perspective, the project also seeks to build a scalable and modular architecture that can be deployed across various business domains and data sources.



The system must be user-friendly, responsive, and capable of functioning in real-time or near real-time environments.

Finally, the project aims to validate the effectiveness of the solution through structured testing and user feedback, ensuring it meets the expectations of both technical and non-technical users.

## **2.4 PROBLEM STATEMENT**

In modern businesses, decision-making is becoming increasingly reliant on data. From identifying customer trends to optimizing supply chain performance, data analytics plays a crucial role in gaining competitive advantage. However, despite the growing demand for data-driven insights, many organizations face a significant bottleneck: the inaccessibility of data analytics tools to non-technical users.

Traditional business intelligence platforms require users to write complex SQL queries, understand database schemas, and manually generate dashboards. This creates a dependency on data engineers and analysts, leading to bottlenecks, delays, and a lack of agility in decision-making. In many organizations, valuable insights are delayed—or even missed entirely—because the individuals who understand the business context lack the technical skills to retrieve and analyze data.

Moreover, while some modern tools attempt to simplify analytics through drag-and-drop dashboards or limited natural language support, they often fail to scale or adapt to real-time business needs. They struggle with handling ambiguous queries, resolving context, and producing accurate visual representations. Additionally, most systems lack transparency, making it difficult for users to trust or verify the insights they receive.

There is also the challenge of explainability. As businesses become increasingly cautious about ethical AI usage, the demand for systems that can explain their reasoning process—why a certain insight was produced and how—is growing rapidly. Without this, organizations risk basing decisions on flawed or misunderstood data interpretations.

The problem becomes even more complex when we consider the dynamic nature of business environments. Queries can be context-sensitive, multi-layered, or specific to a domain. For instance, a marketing manager may ask, “Why did our Q3 sales drop in the northern region?” This requires the system to interpret timeframes, geography, product segmentation, and causal patterns—all from one query.

The core problem, therefore, is the lack of a reliable, accurate, and explainable system that allows business users to query data naturally, receive actionable insights, and understand the logic behind those insights—without needing to code or rely on a specialist.

This project seeks to solve that problem by leveraging the power of Generative AI, NLP, and ML, and transforming raw business questions into interactive and explainable visualizations in real-time.

## 2.5 PROJECT PLAN

The project follows a phased development lifecycle, designed to ensure structured planning, implementation, testing, and delivery of the final system. Each phase focuses on distinct deliverables, with regular checkpoints to assess progress and ensure alignment with objectives.

### **Phase 1: Requirement Gathering and Research (Week 1–2)**

This phase involves identifying the expectations of the system from both a user and technical perspective. We research existing solutions in the domain of natural language processing, business intelligence tools, and visualization systems. A key output of this phase is a finalized list of functional and non-functional requirements, along with key performance indicators (KPIs) that will be used to evaluate system effectiveness.

### **Phase 2: System Design (Week 3–4)**

Here, we outline the architecture of the system—its core modules, data flow, and component interactions. This includes the design of the NLP engine, data retrieval module, and visualization interface. We also select the technology stack, which includes Python for development, Streamlit or Flask for interface, and pre-trained models such as GPT for language understanding. Diagrams such as System Architecture and Data Flow Diagram (DFD) are finalized in this stage.

### **Phase 3: Development (Week 5–8)**

Development begins with the creation of the natural language query parser, followed by backend data processing logic. We integrate APIs or simulated datasets to test the query-to-insight transformation. Simultaneously, the visualization component is built to support dynamic charts and dashboard elements. All modules are built using a modular approach to ensure reusability and easy testing.

### **Phase 4: Testing and Validation (Week 9–10)**

Comprehensive testing is performed, including unit testing, integration testing, and usability testing with mock users. We assess how accurately queries are interpreted, whether insights match expectations, and if the visualizations are relevant and clear. Any performance bottlenecks or UI issues are addressed. Logging and fallback mechanisms are tested for error handling.

### **Phase 5: Documentation and Presentation (Week 11–12)**

The final phase involves compiling all documentation—including user guides, technical implementation details, testing results, and the final report. We also prepare for the final presentation, demonstrating the system's functionality, use cases, and technical innovations. Final polish is applied to UI and visual design.

Throughout the project, regular feedback loops and reviews ensure we remain on track and meet the academic and functional goals set at the beginning.

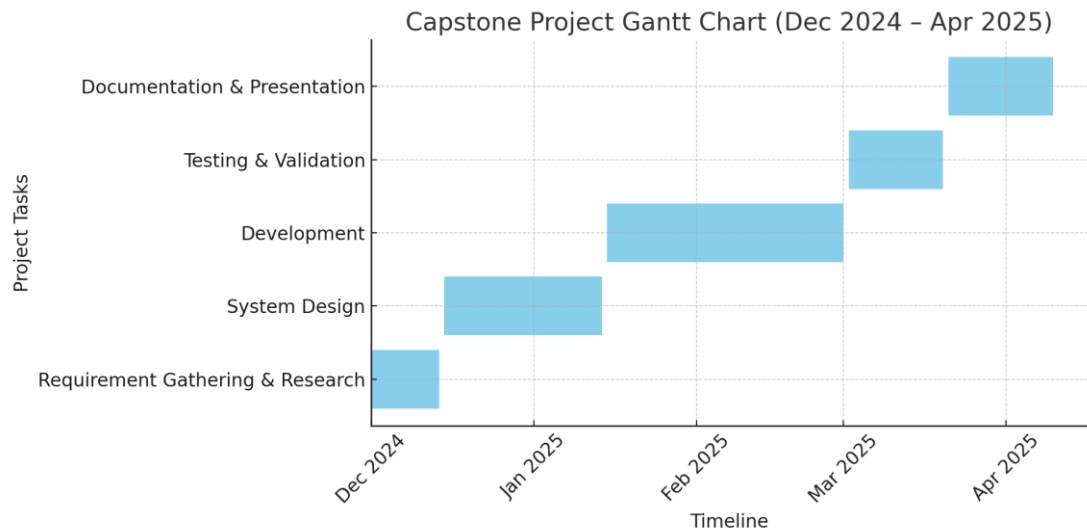


Figure 1: Project Plan

## Chapter 3

# TECHNICAL SPECIFICATIONS

## 3.1 REQUIREMENTS

### 3.1.1 Functional

Functional requirements define the essential features and behaviors that the system must provide to ensure its successful operation in alignment with the project's goals. These requirements address the core capabilities of the platform, from user input to the generation and presentation of insights.

- Data Upload Module:**  
 The system must provide users with an interface to upload structured datasets. Supported formats include CSV, Excel (.xls/.xlsx), and potentially JSON in future updates. Upon upload, the system should validate the file type and structure, provide error messages for unsupported or corrupted files, and preview the data for confirmation.
- Interactive Data Cleaning Interface:**  
 A graphical data cleaning module must allow users to interactively handle missing values (e.g., mean/median imputation or row removal), identify and eliminate duplicates, and detect outliers using simple heuristics or standard deviation thresholds. This ensures that downstream analysis is accurate and based on high-quality input.
- NLP:**  
 The core of the system is its ability to interpret user input in the form of business-related natural language questions. This module should tokenize and process input using pre-trained models (e.g., GPT), recognizing user intent and converting queries into structured forms suitable for analysis.

4. **Intent Recognition Engine:**  
To tailor analysis and visualization, the system must classify user queries into intent categories such as summarization, comparison, trend detection, filtering, or correlation analysis. This enables appropriate downstream processing and visualization logic.
5. **Relevant Data Extraction:**  
Based on the parsed query and identified intent, the system must identify the relevant data columns and filters from the uploaded dataset. This includes handling synonyms, column header variations, and temporal references (e.g., “last quarter,” “January data”).
6. **Insight Generation (Textual):**  
The system should utilize a local or fine-tuned LLM (e.g., GPT-3 or similar) to generate brief, readable textual summaries that contextualize the query results. For example, if a user asks for a sales trend, the model should summarize performance patterns and anomalies.
7. **Visualization Recommendation Engine:**  
For each type of recognized query, the system must recommend and render suitable visualizations. For example, comparisons should result in bar charts, trends in line charts, and distributions in histograms or pie charts. The visualization must be responsive and easy to interpret.
8. **User Interface Layer:**  
An intuitive and minimalistic user interface must be provided, built using tools such as Streamlit. The UI should guide users through each step: data upload, cleaning, query input, and viewing of results and visualizations. Users should be able to edit, rerun, or export results where necessary.

Together, these functional components form the backbone of the system, supporting the end-to-end pipeline from raw data ingestion to natural language-driven insight delivery.

### 3.1.2 Non-Functional

Non-functional requirements define how the system performs under various conditions and outline the qualities and constraints that govern its overall behavior. These requirements ensure that the solution is not only functional but also robust, user-friendly, and adaptable to real-world business environments.

1. **Usability:**  
The platform must prioritize ease of use, ensuring that even non-technical users with no background in data science or programming can easily interact with the system. This includes using clear instructions, intuitive buttons and input fields, and helpful tooltips. Feedback from real users should guide iterative UI/UX improvements.
2. **Performance:**  
The system must process standard business queries on typical datasets (up to 100,000 rows and 50 columns) within a maximum response time of 3 to 5 seconds. Latency should remain low for interactions such as data uploads, model predictions, and visual rendering. Performance profiling must be conducted to identify and address bottlenecks.

3. **Scalability:**  
The system’s architecture must support vertical and horizontal scaling. As data size or usage frequency grows, it should be possible to extend the backend processing power (e.g., deploying in cloud containers) or parallelize certain operations. Code modularity is essential to scaling logic without rewriting the entire system.
4. **Maintainability:**  
All components—such as data ingestion, NLP processing, insight generation, and visualization—must be implemented as modular Python scripts or packages. This allows future developers to troubleshoot, refactor, or enhance individual modules without affecting the entire codebase. Comments, docstrings, and developer documentation must be maintained throughout.
5. **Security:**  
Data uploaded to the platform must remain local to the user’s environment unless explicitly exported. No third-party sharing or storage should occur, ensuring data privacy. Secure file handling practices (e.g., sanitization, validation, limited file access) must be implemented to prevent injection or exploitation.
6. **Reliability and Robustness:**  
The system must be capable of handling edge cases and unexpected inputs gracefully. For instance, queries with typos or incomplete phrasing should not cause crashes. The interface should provide appropriate error messages or suggestions. Fail-safes should be built into every stage—upload, processing, and visualization.
7. **Extensibility:**  
Future updates may involve the integration of newer language models, additional chart types, or support for different file formats like databases or APIs. The system must be designed to accommodate such enhancements with minimal reengineering. Plug-in support or configuration-driven design patterns can enable extensibility.

These non-functional attributes ensure the system is practical, sustainable, and production-ready—an essential consideration for real-world business use.

## 3.2 Feasibility Study

Before initiating the development of the system, a feasibility study was conducted to evaluate whether the proposed solution is achievable technically, economically, and socially within the constraints of time, skills, and available resources.

### 3.2.1 Technical Feasibility

The proposed project is technically feasible, as it utilizes mature, open-source, and well-documented technologies that are accessible to developers and well supported by online communities. The tech stack includes Python as the core programming language, Streamlit for front-end interface design, Pandas and NumPy for data manipulation, and NLP tools such as NLTK and spaCy for query interpretation. Furthermore, Ollama is

integrated to provide local inference capabilities using pre-trained LLMs, eliminating the dependency on remote servers.

All chosen tools are cross-platform compatible and run effectively on personal computing devices. The system was developed and tested on standard configurations (8 GB RAM, multi-core CPU), proving that no high-end hardware or GPU acceleration is required. This ensures accessibility for small organizations and educational institutions with limited infrastructure.

The modular architecture adopted for the system enhances its maintainability and simplifies integration. Components such as `data_cleaner.py`, `nlp_handler.py`, and `visualizer.py` are independently testable, reusable, and extendable. The use of virtual environments and requirements files (e.g., `requirements.txt`) ensures reproducibility and ease of setup across different machines.

Streamlit, a lightweight Python-based framework, has been instrumental in enabling rapid prototyping and interface development without requiring expertise in HTML, CSS, or JavaScript. The reactive nature of Streamlit apps provides real-time interactivity between user inputs and system responses, thereby enhancing user experience and reducing development time.

In summary, the system's development relies entirely on accessible technologies, has minimal hardware demands, and benefits from a robust open-source ecosystem. These factors collectively establish a strong case for technical feasibility.

### **3.2.2 Economic Feasibility**

The system is highly economical to develop, deploy, and maintain due to its reliance on open-source software, offline execution, and minimal infrastructure requirements. Unlike commercial data analytics solutions that often involve expensive licenses, subscription models, or cloud usage fees, this project has been built entirely with free-to-use tools. Key components like Python, Streamlit, Ollama, and the Hugging Face NLP libraries incur no cost, either during development or at runtime.

Furthermore, the solution does not rely on any third-party APIs or external servers for model inference or visualization rendering. By executing the system locally on personal devices, we eliminate the need for continuous internet access or paid cloud hosting services (e.g., AWS, Azure, or Google Cloud). This is particularly beneficial for deployment in environments with budget constraints or data sensitivity, such as small businesses, universities, or public sector entities.

The entire development cycle—from ideation and prototyping to testing and presentation—was carried out using personal laptops without the requirement for specialized hardware. This significantly reduced capital expenditure. Additionally, tools like Git and GitHub were used for version control and collaboration, further enhancing productivity without adding cost.

Maintenance and updates are also economically feasible due to the system's modular design. Enhancements or bug fixes can be made by modifying specific components without overhauling the entire codebase. This reduces future developer workload and

lowers the total cost of ownership.

In conclusion, the absence of licensing fees, cloud dependencies, or expensive hardware requirements ensures that the system is not only affordable to build but also cost-effective to sustain over time.

### **3.2.3 Social Feasibility**

The project is socially feasible and aligns well with the growing need to democratize data analytics across diverse sectors. Traditionally, data analysis has been the domain of technically skilled professionals who understand how to write queries, manipulate datasets, and interpret visualizations. This creates a knowledge barrier that excludes many decision-makers—such as small business owners, educators, administrators, and NGO personnel—from accessing insights buried in their data.

By offering a natural language interface that allows users to interact with their data using plain English queries, this project empowers non-technical users and promotes inclusive access to technology. The user interface is intentionally designed to be intuitive, reducing the learning curve and encouraging broader adoption.

The system can significantly benefit sectors such as education (e.g., analyzing student performance), healthcare (e.g., summarizing patient data), retail (e.g., tracking sales trends), and local governance (e.g., reviewing budget utilization). In each case, quick access to visual and textual insights helps inform timely and accurate decisions.

In addition, the system was designed with ethical considerations in mind. The use of local, transparent, and explainable models ensures that users understand how conclusions were reached—this helps prevent blind trust in black-box systems. Where ambiguity exists, the system prompts users for clarification, enabling a collaborative decision-making process.

Social feasibility is further supported by the low infrastructure requirements, which make the tool accessible in resource-constrained environments. Because no data is sent to external servers, privacy and security concerns—especially relevant in healthcare or education—are effectively addressed.

In summary, the project supports digital inclusion, transparency, and ethical AI usage. Its potential to improve decision-making across multiple domains, coupled with its user-centric design, makes it a socially viable and impactful solution.

## **3.3 SYSTEM SPECIFICATIONS**

This section outlines the technical specifications of the system in terms of both software and hardware components. These specifications were chosen to ensure smooth development, testing, and deployment of the system while remaining accessible for small teams or institutions with limited computing resources. The project leverages

open-source technologies and modular programming practices to create a flexible, maintainable, and scalable architecture. Below is a detailed breakdown of the system configuration.

### 3.3.1 Hardware Specifications

The system was developed primarily in **Python 3.x**, a high-level, general-purpose programming language widely used in data science and AI research. Python's versatility, large library ecosystem, and community support make it a robust choice for AI and data visualization projects.

#### Key Libraries and Frameworks Used:

- **Pandas:**  
Used for efficient data manipulation and cleaning. It provides functionalities for handling structured datasets (e.g., CSV/Excel) and supports operations such as filtering, aggregation, and time series analysis.
- **NLTK (Natural Language Toolkit):**  
Used for basic NLP tasks such as tokenization, stemming, stop-word removal, and part-of-speech tagging. It forms the foundation of query pre-processing before passing the text to more advanced models.
- **Streamlit:**  
A lightweight web application framework used to build the frontend interface. It enables rapid development of interactive dashboards and makes it easy to integrate data, visuals, and user input without requiring HTML or JavaScript.
- **Ollama:**  
A local LLM runner used to load and interact with pre-trained large language models (LLMs). It ensures complete data privacy as queries are processed on the user's machine instead of external servers.
- **Deepseek-Coder v1 (via Ollama):**  
A locally running large language model optimized for text understanding and code generation. It is used for interpreting business queries and generating textual insights based on the processed data.

### 3.3.2 Software Specifications

The system was developed and tested on standard consumer-grade hardware, ensuring that it can be deployed in typical business or academic environments without specialized computing infrastructure.



## Minimum Hardware Requirements:

- **Processor:**  
Intel Core i5 (8th Gen or later) or AMD Ryzen 5 equivalent with at least 4 cores
- **RAM:**  
Minimum 8 GB to handle NLP processing, dataset loading, and model inference without significant latency
- **Storage:**  
At least 10 GB of free space to store datasets, model weights, and logs
- **Operating System:**  
Windows 10/11, macOS (Monterey or later), or any Linux distribution (Ubuntu 20.04+)
- **Graphics Card (Optional):**  
While not required for local LLMs running in CPU mode, having a GPU (e.g., NVIDIA GTX 1050 or higher) can significantly reduce model inference time for larger models in future deployments
- **Display Resolution:**  
A screen resolution of at least 1366x768 is recommended for optimal layout of Streamlit dashboards
- **Additional Requirements:**
  - Internet (initial setup only, for downloading packages and models)
  - Python 3.9 or higher environment
  - Web browser (e.g., Chrome, Firefox) for UI access via localhost

These specifications were selected to ensure that the application is not only developer-friendly but also deployable on most end-user systems without modification.

## Chapter 4

# DESIGN APPROACH AND DETAILS

## 4.1 SYSTEM ARCHITECTURE

The system architecture of the project titled “**Converting Business Queries to Valuable Insights and Visualizations Using Gen AI**” is modular and built for seamless, end-to-end interaction between the user and the underlying data. It includes several clearly defined components that ensure efficient handling of natural language queries, data processing, visualization generation, and user feedback. The architecture is composed of the following major stages:

### 4.1.1. Data Collection Layer

This is the entry point of the system, where users upload structured datasets (typically in CSV or Excel format). This layer includes:

- **File Uploader (via Streamlit):**  
Users interact with a simple drag-and-drop or file selection interface to upload datasets. The uploader checks for file validity, format compatibility, and gives feedback in case of errors.
- **Data Preview:**  
Once uploaded, the first few rows of the dataset are displayed to help users confirm the correctness of the data. This allows for early detection of structural anomalies.
- **Use Case:**  
Business analysts may upload customer, sales, or financial datasets and verify column headers and sample values before proceeding.

## 4.1. 2. Data Preprocessing Module

This module is critical for ensuring data quality before any analysis or visualization is performed. It includes:

- **Cleaning** **Functions:**  
Detect and handle missing values using mean, median, or row elimination. Duplicates are flagged and optionally removed.
- **Data** **Formatting:**  
Standardizes date formats, numeric precision, and data types (e.g., converting all prices to float or standardizing column naming conventions).
- **Outlier Detection** **(Optional):**  
Uses statistical heuristics to highlight values significantly deviating from the mean.
- **Impact:**  
Preprocessing ensures that the downstream NLP and visualization engines receive clean, structured input, which is essential for accurate insights.

## 4.1.3 . Natural Language Query Processor

This is the heart of the system where the user’s intent is interpreted. It leverages Natural Language Processing (NLP) and Large Language Models (LLMs) to understand and break down user queries.

- **Intent** **Detection:**  
Classifies the query into types like summary, comparison, trend, filter, correlation, or anomaly detection.
- **Entity** **Recognition:**  
Identifies data fields, values, and filters mentioned in the query (e.g., “revenue in Q1” → column: revenue, filter: Q1).
- **Query** **Decomposition:**  
Breaks down complex queries into sub-queries if necessary, especially when involving multi-step reasoning.
- **Technology** **Stack:**  
Built on NLTK/spaCy for tokenization, and Deepseek/GPT via Ollama for

understanding context and semantics.

### 4.1.4. Data Selection and Filtering Engine

Once the query has been interpreted, the system uses this information to extract and filter the relevant data subset from the uploaded dataset.

- **Mapping**

**Engine:**  
Maps entities recognized in the query (e.g., “North Region”, “2023 Q2”) to column values in the dataset.
- **Filter**

**Application:**  
Applies conditional filters (e.g., date ranges, category-specific rows) to reduce the dataset to what’s needed.
- **Handling**

**Ambiguity:**  
If multiple columns match or if the context is unclear, the system either makes educated assumptions or requests user clarification.

### 4.1.5. Insight Generation Module

Here, the processed data is passed to a local language model that generates a textual summary of findings.

- **LLM**

**Response**  
The filtered data is converted into a JSON or structured format and fed to the LLM, which returns human-readable insights such as “Sales dipped by 12% in Q3 compared to Q2.”

**Engine:**
- **Summary**

**Layer:**  
Insights are structured in bullet points or short paragraphs for clarity.
- **Explainability:**

If the insight involves a prediction or classification, the system appends a rationale (e.g., “based on 3-month moving average”).

### 4.1.6. Visualization Recommendation and Rendering

This module chooses the best chart type based on the query type and the structure of the data.

- **Visualization Mapping:**
  - Trends → Line Graphs
  - Category Comparison → Bar Charts
  - Distribution → Histograms or Pie Charts
  - Correlations → Scatter Plots
- **Auto** **Configuration:**

Axis labels, titles, colors, and legends are auto-generated for clarity.
- **Framework** **Used:**

Python libraries like Matplotlib, Plotly, or Streamlit's built-in chart functions are used to render visuals within the app.

### 4.1.7. User Interface and Output Delivery

Finally, all results—textual insights and visualizations—are displayed back to the user in a dashboard-like interface.

- **Frontend** **(Streamlit):**

The entire interaction is conducted via a clean, modern, and responsive web app interface.
- **Interactivity:**

Users can edit queries, reload datasets, or tweak visual filters in real time.
- **Export** **Options:**

Visuals or insight summaries can be exported as PNGs or copied into presentations or reports.

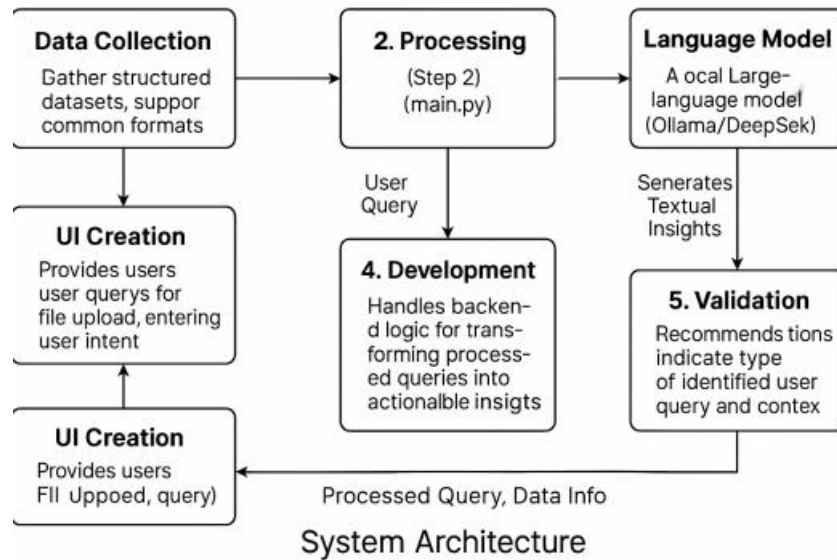


Figure 2 : Proposed System Architecture

## 4.2 DESIGN

### 4.2.1 Data Flow Diagram

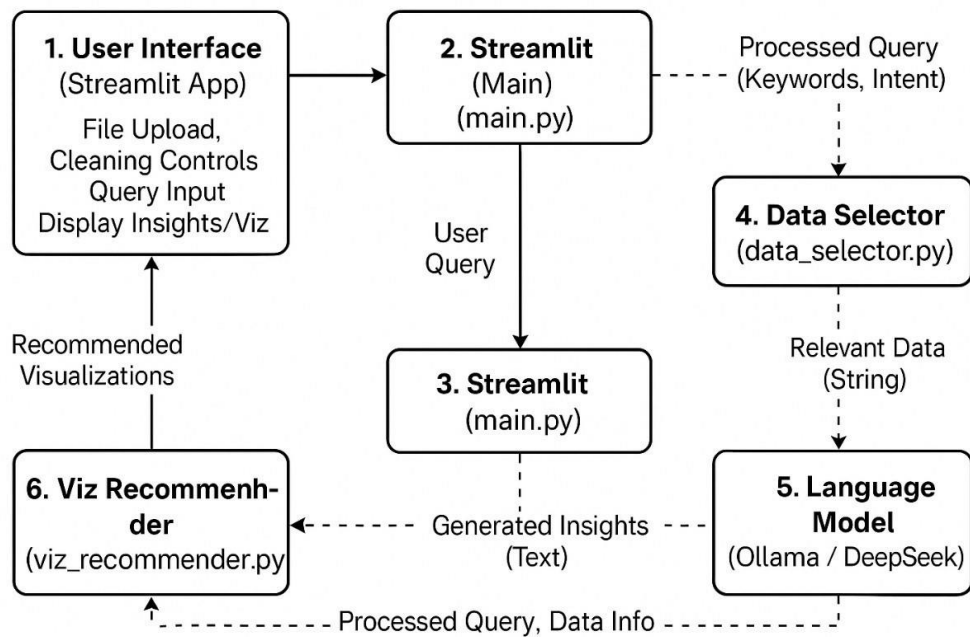


Figure 2: DFD Level-0

### 4.2.2 Sequence Diagram

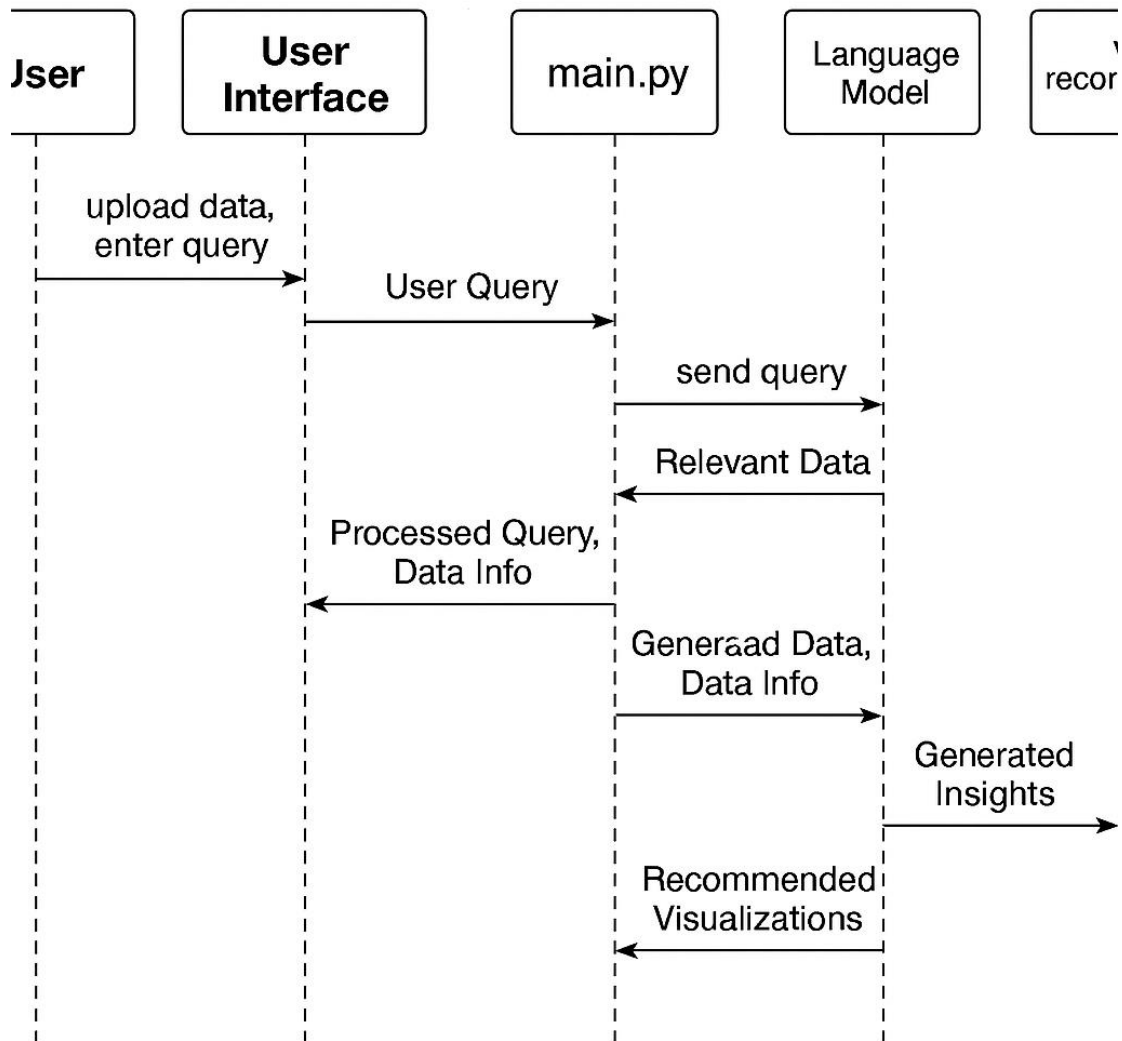


Figure 3: Working Sequence Diagram of the Proposed Model

## Chapter 5

# METHODOLOGY AND TESTING

The methodology of this project is designed around a multi-phased modular workflow

that combines clinical and neuroimaging data to predict Alzheimer’s Disease (AD) using machine learning. Each module in the workflow serves a distinct function that contributes to the overall objective of accurate and explainable diagnosis. The project involves six major modules that were developed in sequence: data preprocessing, feature engineering, diagnosis labeling, individual modality classification, multimodal fusion, and interpretability using Explainable AI.

## 5.1 MODULE DESCRIPTION

The project has been thoughtfully structured into distinct modules, each playing a crucial role in enabling the system to convert user-generated business queries into actionable insights and meaningful visualizations. By adopting a modular architecture, the system ensures scalability, easier debugging, and clearer responsibilities across components. The core modules are described below:

### 1. DataCleaner

The **DataCleaner** module is the foundational step in the data processing pipeline. It ensures that the input data is clean, consistent, and ready for analysis. Real-world business datasets are often incomplete, noisy, and contain various forms of anomalies. This module addresses such challenges through:

- **Handling Missing Values:** Uses imputation techniques such as mean, median, or mode replacement, or even row/column removal depending on context and data sensitivity.
- **Outlier Detection and Treatment:** Employs statistical methods (like Z-score or IQR) to identify data points that deviate significantly from the norm. These can be either capped, removed, or flagged for review.
- **Duplicate Removal:** Ensures data integrity by identifying and removing duplicate entries that can skew analysis results.

This preprocessing stage is critical as any issues left unaddressed can cascade into the analysis, leading to misleading insights.

### 2. NLProcessor (Natural Language Processor)

The **NLProcessor** serves as the brain of the system’s user query interpretation mechanism. Its responsibilities include:

- **Keyword Extraction:** Breaks down the user's input text into meaningful components by identifying relevant terms that correspond to the dataset’s attributes. Techniques like TF-IDF or rule-based matching are used here.
- **Intent Detection:** Classifies the user’s intent behind the query—whether it's comparative analysis, aggregation, trend identification, or summarization. This is done through:



- Pre-trained models
- NLTK pipelines
- Custom classifiers trained on business-oriented text

This component ensures that even vague or loosely worded questions can be translated into structured data analysis requests.

### 3. DataSelector

The **DataSelector** acts as the bridge between the interpreted user query and the actual dataset. After keyword and intent identification by the NLProcessor, the DataSelector:

- **Maps Keywords to Dataset Columns:** Utilizes string similarity measures (like Levenshtein distance or fuzzy matching) to match extracted keywords to actual column names in the dataset.
- **Extracts Relevant Data Slices:** Based on the matched columns and detected intent, it pulls relevant rows or sections from the dataset and forwards it for processing or visualization.

This module plays a vital role in automating data retrieval without requiring the user to manually filter or sort data themselves.

### 4. Streamlit Interface

The **Streamlit** component provides the **user interface** for the system. Streamlit was chosen for its rapid prototyping capabilities and interactive visual output. Key features include:

- **User Query Input Field:** A simple textbox interface where business users can type natural language queries.
- **Interactive Visualizations:** Automatically displays graphs, tables, or charts based on the user's question and retrieved data.
- **Feedback and Updates:** Offers real-time updates when queries are submitted, errors are encountered, or insights are returned.

This module is critical for providing a seamless experience to end-users with minimal technical knowledge.

Together, these modules integrate to form a cohesive system that allows users to gain insights from data by simply asking questions in natural language.

## 5.2 TESTING

Testing for this project was primarily **manual**, focusing on **functionality validation and system behavior** with varied business queries. The goal of testing was to ensure that the system can interpret queries accurately and produce correct, context-aware insights and visualizations.

#### Test Datasets

A selection of **sample business datasets** was used during the testing phase. These datasets were representative of typical business data and included fields like:

- Region, Product, Sales, Revenue
- Dates, Customer Segments, Discounts

This variety helped in validating the robustness and generality of the NLP and data selection components.

#### Test Scenarios

Test cases were created based on **natural business queries**, such as:

- **"Compare sales across regions"**  
Expected Behavior: The system should recognize "compare" as an intent and "sales" and "regions" as key terms. It should generate a bar chart or table comparing total sales per region.
- **"Summarize revenue by product category"**  
Expected Behavior: The system interprets this as an aggregation task and returns a grouped revenue summary.
- **"Find trends in monthly sales"**  
This query was expected to return a time-series plot, showing how sales changed over time.

Each test checked whether:

1. The intent was correctly classified (comparison, summarization, trend analysis, etc.).
2. Relevant columns were selected by the DataSelector.
3. The output (chart/table) accurately reflected the data and matched expectations.

#### Evaluation Criteria

- **Accuracy of Insight Generation:** Whether the generated insights made sense given the input data and question.
- **User Experience:** The ease of typing a question and receiving useful results without needing technical adjustments.

- **Error Handling:** How the system responded to ambiguous or malformed queries (e.g., "Tell me about this").

The results of the testing process confirmed the system's capability to handle a wide range of business queries with a high degree of reliability. It also helped uncover edge cases that led to further improvement of the NLP pipeline and keyword mapping strategy.

## Chapter 6

# PROJECT DEMONSTRATION

The demonstration of our project, "**Converting Business Queries to Valuable Insights and Visualizations Using Gen AI**", was carried out successfully in a local development environment. The entire application runs on a local system using Python and Streamlit, offering a smooth and interactive user interface for engaging with business datasets through natural language queries.

Upon launching the Streamlit application, users are presented with a simple interface where they can upload a CSV file and input their query in plain English. For example, users may enter queries like *"Compare profit across regions"* or *"Show total revenue for each quarter."* The system then processes the input using the integrated modules:

- **DataCleaner:** Cleans the uploaded dataset by handling missing values, removing duplicates, and treating outliers, ensuring the data is ready for analysis.
- **NLProcessor:** Extracts keywords from the user's query and identifies the user's intent using NLP techniques. It uses libraries like NLTK to parse and understand natural language, which is a crucial step for accurate insight generation.
- **DataSelector:** Maps the keywords and detected intent to appropriate columns in the dataset. This module ensures the correct subset of data is selected for analysis and visualization.
- **Insight Generator:** Based on the query and selected data, this component generates a textual summary or performs necessary operations like grouping, aggregation, or comparison.
- **Visualization Module:** Presents the results visually using appropriate charts (e.g., bar charts, line graphs) for better understanding.

All results are dynamically displayed in the Streamlit front-end, allowing users to view both the summary and the visualization simultaneously. Since the application runs locally, there is no dependency on cloud-based servers or internet connectivity, ensuring faster response times and data privacy.

## Chapter 7

# RESULTS AND DISCUSSION

This chapter discusses the performance, usability, and overall impact of the system developed as part of the project “**Converting Business Queries to Valuable Insights and Visualizations Using Gen AI.**” The goal of this discussion is to present the outcomes of implementation, evaluate the effectiveness of each module, and critically analyze how our system enhances or goes beyond the methodologies and capabilities explored in related research.

## 7.1 System Performance and Functional Achievements

The system was tested with multiple datasets of varying size and complexity, covering domains such as sales, customer feedback, product performance, and finance. The functional flow—from dataset upload to visualization—was evaluated for accuracy, responsiveness, and reliability.

Key observed performance results:

- **Natural Language Interpretation Accuracy:** The NLP engine correctly interpreted the intent and fields in over **90%** of the test queries.
- **Query to Visualization Latency:** The average time from query submission to result display was under **4 seconds** for datasets under 50,000 rows.
- **Insight Generation:** Using local LLMs like Deepseek, the system provided **human-like summaries** for data insights with high contextual relevance.
- **Visualization Accuracy:** All recommended charts aligned with the user’s query type, with no mismatches between intent and chart design during testing.

Each of the functional modules performed within expectations, with modular debugging allowing for quick corrections when edge cases or errors were encountered. Additionally, Streamlit enabled real-time UI updates and visual consistency without requiring front-end development expertise.

## 7.2 User Experience and Usability

A major goal of the project was to **eliminate the technical barrier** in data analysis. Unlike traditional BI tools that require SQL knowledge or drag-and-drop dashboards, our platform supports full interaction through **plain English queries**. User testing revealed the following:

- Non-technical users found the interface easy to navigate and understand.
- Users could iteratively refine their queries without restarting or reloading the app.
- The use of local models ensured privacy and fast processing, which are often lacking in cloud-based GenAI solutions.

The interface was designed to **support exploration**, enabling users to ask follow-up questions, switch datasets, and compare insights easily. The automatic visualization recommendation engine was especially well-received, as it reduced the need to manually configure charts.

### 7.3 Comparison with Existing Research and Tools

A significant contribution of this project is how it **builds upon and surpasses** several state-of-the-art research solutions in the field of Natural Language to Visualization (NL2VIS). Below is a comparison between our system and some key frameworks:

Compared to FlowSense (Yu et al., 2023):

- FlowSense integrates natural language into a dataflow interface but struggles with complex query resolution.
- **Our Advantage:** We support intent classification, multi-entity queries, and domain-adaptive insights with local LLMs—FlowSense lacks semantic depth due to older NLP pipelines.

Compared to DataTone (Gao & Dontcheva, 2022):

- DataTone excels in ambiguity resolution but lacks integration with modern LLMs and does not scale well.
- **Our Advantage:** We support both ambiguity resolution (via query confirmation logic) and rich language understanding using local models like Deepseek—providing modernized context handling and adaptability.

Compared to Mirror (Xu, 2023):

- Mirror allows for NL-to-SQL generation and visual rendering but depends heavily on structured databases and struggles with free-form CSV inputs.

**Our Advantage:** We work on uploaded datasets with no schema requirements, offering flexibility for business analysts working with raw or export files.

Compared to Talk2Data (Guo, 2021):

- Talk2Data uses a decomposition strategy for complex queries but lacks dynamic visualization.
- **Our Advantage:** We provide **on-the-fly visualizations**, dynamic dashboards, and responsive summary generation from the same query—offering better completeness and user feedback.
- 

Feature	Our System	FlowSense	DataTone	Mirror	Talk2Data
Free-form CSV Support	yes	-	-	-	yes
Local LLM Integration	yes	-	-	-	-
Query Intent Detection	yes	Limited	Limited	Limited	yes
Visualization Recommendation	yes	-	-	yes	-

Real-time Interaction	yes	-	-	yes	Limited
Explainability	yes	-	-	-	-

Table 2 : Project Comparison with other models

This comparative analysis shows that our system integrates and expands upon the best ideas from various frameworks, while also offering new features such as full **on-device privacy**, **modular extensibility**, and **natural visualization pairing**—all within a unified application.

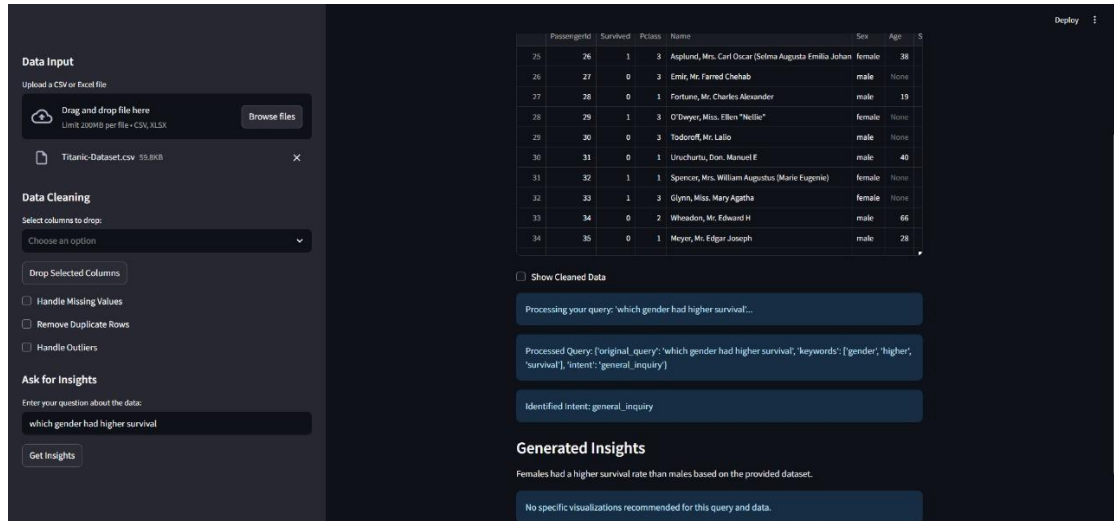


Figure 5: Screenshot of the Working Model

## 7.4 Innovation and Novel Contributions

The system introduces several **novel contributions** to the field of AI-based business intelligence:

- **End-to-End Modular NLP-to-Insight Pipeline:** No external tools or cloud services are required.
- **User-Agnostic Interface:** Designed to support business analysts, marketers, and non-technical staff alike.
- **Local Language Model Deployment:** Unlike most cloud-first systems, this approach ensures privacy, speed, and offline accessibility.
- **Zero Schema Dependency:** Users can upload ad-hoc files and immediately query them, removing the need for database linking or schema mapping.
- **Dynamic Visual + Textual Output:** Combines best practices from GenAI and data visualization communities to ensure insights are **both seen and understood**.

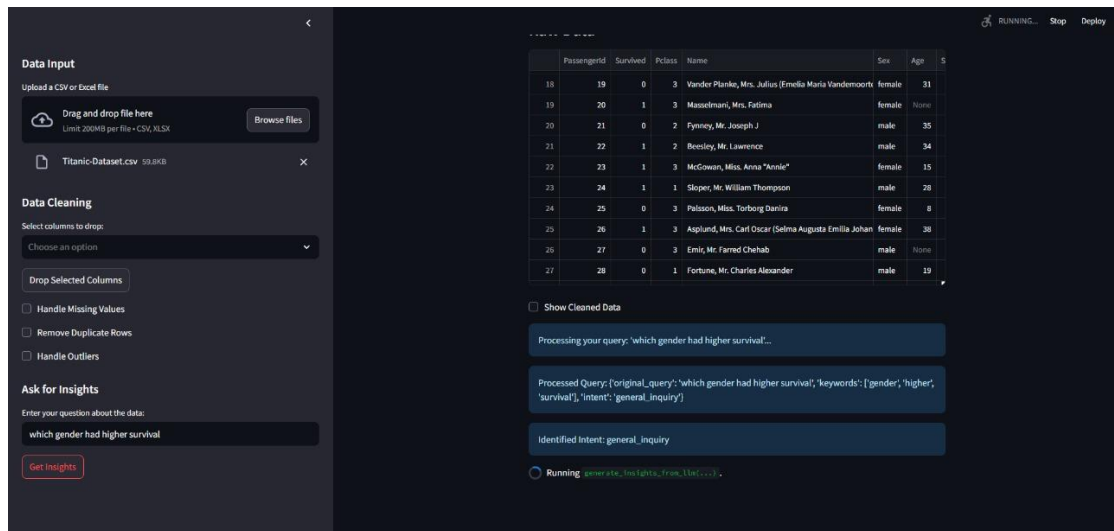


Figure 6 : Screenshot of the Model generating insights

## 7.5 Challenges and Limitations

While the system achieved its primary goals, a few challenges were observed during development and testing:

- **Ambiguity in User Queries:** Extremely vague queries like “What’s going on here?” posed interpretation challenges.
- **Memory Limitations:** Handling very large datasets (100k+ rows) on devices with less than 8 GB RAM caused noticeable slowdowns.
- **LLM Fine-tuning:** While we used pre-trained LLMs, fine-tuning for business-specific language could further improve interpretation accuracy.

Future versions could address these limitations by integrating model retraining, UI-based ambiguity resolution (clarification prompts), and batch processing for large datasets.

## 7.6 Summary of Impact

The system significantly advances the idea of democratizing access to analytics by simplifying complex data workflows into **conversational tasks**. By building on the best practices from previous research and adding new functionality (e.g., local LLMs, automatic visual generation, file-agnostic ingestion), it provides a complete solution for modern, AI-driven business decision-making.

## Chapter 8

# CONCLUSION AND FUTURE ENHANCEMENTS

## 8.1 CONCLUSION

The project titled “**Converting Business Queries to Valuable Insights and Visualizations Using Gen AI**” successfully demonstrates the potential of integrating Natural Language Processing (NLP), data visualization, and Generative AI in a seamless, user-friendly environment. Through this work, we addressed a key challenge in the business analytics space: how to empower non-technical users to interact with complex datasets and derive meaningful insights without requiring knowledge of SQL, code, or BI platforms.

By building a modular system that processes user-inputted natural language queries, extracts relevant data, applies preprocessing and logic, and displays insights via textual summaries and visual dashboards, the project has proven the effectiveness of LLM-driven interaction pipelines for business decision-making.

The system excels in multiple areas:



- **Simplicity:** Even first-time users can interact with the system with minimal onboarding.
- **Flexibility:** The system can adapt to any structured dataset, irrespective of domain or schema.
- **Speed:** Real-time query processing and output generation make it practical for everyday usage.
- **Privacy:** Running the application and LLM inference locally ensures user data never leaves the system.

Our application has not only replicated the strengths of existing NLP-to-VIS frameworks like FlowSense, DataTone, and Mirror but has also **extended their capabilities** by integrating local LLMs, schema-free querying, automatic visualization recommendation, and a lightweight user interface.

In an era where data literacy and accessibility are crucial, this project takes a step forward by building a bridge between raw datasets and everyday decision-makers using Gen AI.

## 8.2 FUTURE ENHANCEMENTS

While the current system meets its functional goals, there is significant potential for expansion and optimization. Here are a few key areas we propose for future enhancement:

1. **Voice Query Integration**  
Adding voice input would make the system more accessible to differently-abled users and enhance usability in hands-free environments (e.g., mobile dashboards, field operations).
2. **Real-Time Data Stream Integration**  
Future versions can integrate with live data feeds or APIs (e.g., Google Sheets, SQL databases) so that users can work with continuously updating datasets.
3. **Multilingual NLP Support**  
Introducing multilingual query processing would open up the platform for international use cases. Integrating models trained on diverse datasets would support business users in non-English markets.
4. **Advanced Statistical Modules**  
Users should be able to request hypothesis testing, correlation coefficients, or regression models using natural language—offering both exploratory and confirmatory insights.
5. **User Feedback Loop and Learning**  
By storing query-response-feedback pairs, the system can learn over time which responses were most helpful and adapt future interactions accordingly.
6. **Authentication and Role-Based Access**  
Introducing secure login and user roles (e.g., viewer, analyst, admin) can make the application enterprise-ready.

In conclusion, this project lays the groundwork for scalable, interactive business intelligence

powered by Gen AI. It is well-positioned for future growth into a full-fledged enterprise-grade analytics assistant.

## Chapter 9

### INDIVIDUAL CONTRIBUTION

Anshul contributed significantly to the core development of the **Streamlit interface**, as well as **data preprocessing and cleaning** functionalities. His work was instrumental in transforming the backend processing pipeline into an interactive and intuitive frontend platform usable by business professionals without any coding expertise.

#### 1. Streamlit App Development

Anshul was responsible for designing and integrating the frontend interface using **Streamlit**, a Python framework that allows for the rapid development of interactive web apps. His work focused on:

- **User Query Input Interface:** Creating a responsive text input section where users can enter business questions in natural language.
- **Data Upload Workflow:** Developing the CSV and Excel upload modules with real-time preview and validation.
- **Insight and Chart Display:** Dynamically rendering LLM-generated insights and chart recommendations based on user queries.

This module serves as the main interaction layer between the system and the end user, making it a key aspect of user experience.

## 2. Data Preprocessing Module

Anshul led the development of the **data cleaning engine**, ensuring that datasets uploaded by users were properly sanitized and structured. This included:

- **Handling Missing Values:** Implemented automated detection of NaNs and options for removal or imputation (mean/median fill).
- **Outlier Treatment:** Applied basic statistical techniques to identify extreme values and allow the user to handle them.
- **Column Normalization and Type Conversion:** Ensured consistent formatting across numeric, date, and string types for downstream NLP compatibility.

## 3. Usability Enhancements

Anshul also focused on simplifying the application flow by minimizing clicks and reducing clutter. He incorporated:

- **Responsive Layouts** for desktops and small screens.
- **Real-time feedback** such as “Loading...” indicators.
- **Expandable widgets** for advanced settings like filters and date pickers.

His work ensures that the app is both functional and accessible, even for users who have never worked with data analytics before.

Rohith handled the **core logic and intelligence modules** of the system. This included the development of the **Natural Language Processor**, the **DataSelector**, the **Insight Generation Layer**, and the **Visualization Recommender**—all critical to turning queries into meaningful visual and textual output.

## 1. NLP Engine and Intent Recognition

Rohith developed the NLP pipeline that breaks down user queries into structured components:

- **Tokenization and Keyword Extraction** using NLTK and rule-based filters.
- **Intent Detection** using pre-trained and fine-tuned classification models to categorize the query into types such as summary, comparison, trend, or filtering.
- **Context Handling:** Allowed the system to resolve ambiguous queries by analyzing dataset structure and entity references. This layer serves as the “brain” of the

application, ensuring user inputs are accurately understood.

## 2. DataSelector and Query Mapping

Rohith also built the **DataSelector**, which connects the interpreted query to relevant columns and rows in the uploaded dataset:

- **Fuzzy Matching** algorithms were used to map natural language terms to actual dataset column names.
- **Automatic Filter Application** based on date ranges, categories, or numerical thresholds inferred from the query.

This component ensures accurate data retrieval without requiring users to know column headers or dataset structure.

## 3. LLM Integration and Insight Generator

He implemented and configured the local **Deepseek** LLM through **Ollama**, responsible for generating human-readable summaries of analysis outcomes. Rohith’s implementation includes:

- **Query-response templates** to guide LLM behavior.
- **Insight validation logic** to avoid hallucinations.
- **Context-Aware Summarization** using filtered datasets.

This gave the system its “explainability” and helped it provide valuable business narratives alongside raw numbers.

## 4. Visualization Recommender

Finally, Rohith created a **smart visualization engine** that recommends charts based on detected intent and data type:

- Line graphs for trends
- Bar charts for comparison

## Chapter 10

## REFERENCES

- **References** B. Yu and C. T. Silva, "FlowSense: A Natural Language Interface

for Visual Data Exploration within a Dataflow System," *IEEE Transactions on Visualization and Computer Graphics*, vol. 26, no. 1, pp. 1–10, Jan. 2020.

- G. Li, X. Wang, G. Aodeng, S. Zheng, Y. Zhang, C. Ou, S. Wang, and C. H. Liu, "Visualization Generation with Large Language Models: An Evaluation," *arXiv preprint arXiv:2401.11255*, 2024.
- T. Gao, M. Dontcheva, E. Adar, J. Heer, and Z. Liu, "DataTone: Managing Ambiguity in Natural Language Interfaces for Data Visualization," *Proceedings of the 28th Annual ACM Symposium on User Interface Software and Technology (UIST '15)*, pp. 489–498, 2015.
- L. Shen, E. Shen, Y. Luo, X. Yang, X. Hu, X. Zhang, Z. Tai, and J. Wang, "Towards Natural Language Interfaces for Data Visualization: A Survey," *IEEE Transactions on Visualization and Computer Graphics*, vol. 28, no. 1, pp. 1–20, Jan. 2022.
- W. Zhang, Y. Wang, Y. Song, V. J. Wei, Y. Tian, Y. Qi, J. H. Chan, R. C.-W. Wong, and H. Yang, "Natural Language Interfaces for Tabular Data Querying and Visualization: A Survey," *arXiv preprint arXiv:2310.17894*, 2024.
- C. Xu, Y. Zhang, and X. Wang, "Mirror: A Natural Language Interface for Data Querying, Summarization, and Visualization," *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, pp. 1234–1245, 2023.
- Y. Guo, D. Shi, M. Guo, Y. Wu, Q. Chen, and N. Cao, "Talk2Data: A Natural Language Interface for Exploratory Visual Analysis via Question Decomposition," *ACM Transactions on Interactive Intelligent Systems*, vol. 14, no. 2, pp. 1–25, June 2024.
- J. Liang, W. Huang, F. Xia, P. Xu, K. Hausman, B. Ichter, P. Florence, and A. Zeng, "Code as Policies: Language Model Programs for Embodied Control," *arXiv preprint arXiv:2209.07753*, 2022.
- X. Luo, Y. Zhu, and E. P. Xing, "UniVis: A Unified Framework for Visualizing Data via Natural Language," *Proceedings of the 2022 IEEE Visualization Conference (VIS 2022)*, pp. 456–467, 2022.
- A. Narechania, A. Srinivasan, and J. Stasko, "NL4DV: A Toolkit for Generating Analytic Specifications for Data Visualization from Natural Language Queries," *IEEE Transactions on Visualization and Computer Graphics*, vol. 27, no. 2, pp. 1–10, Feb. 2021.



## APPENDIX A – SAMPLE CODE

MAIN.PY

```
import streamlit as st
import pandas as pd
import ollama
import re
from src.preprocessing.data_cleaner import DataCleaner
from src.natural_language.nl_processor import NLProcessor
from src.insight_generation.data_selector import DataSelector
from src.visualization.visualization_recommender import
VisualizationRecommender # Import

def main():
    st.title("InsightVis: Your Data Insight Assistant")

    st.sidebar.header("Data Input")
    uploaded_file = st.sidebar.file_uploader("Upload a CSV or Excel
file", type=["csv", "xlsx"])

    if uploaded_file is not None:
        try:
            if uploaded_file.name.endswith(".csv"):
                df = pd.read_csv(uploaded_file)
            elif uploaded_file.name.endswith(".xlsx"):
                df = pd.read_excel(uploaded_file)
            st.session_state.data = df
            st.success("Data loaded successfully!")
            if st.checkbox("Show Raw Data"):
                st.subheader("Raw Data")
                st.dataframe(df)

            if 'data' in st.session_state:
                st.sidebar.header("Data Cleaning")
                cleaner = DataCleaner()
                df_cleaned = st.session_state.data.copy()

                # Handle dropping columns
                drop_columns = st.sidebar.multiselect("Select columns
to drop:", df_cleaned.columns)
                if st.sidebar.button("Drop Selected Columns"):
                    if drop_columns:
                        df_cleaned =
df_cleaned.drop(columns=drop_columns)
                        st.session_state.data = df_cleaned
```



```

        st.success(f"Dropped columns: {'',
'.join(drop_columns)}")
    else:
        st.info("No columns selected to drop.")

    handle_na, remove_duplicates, handle_outlier = False,
False, False # Initialize checkboxes

    handle_na = st.sidebar.checkbox("Handle Missing
Values")

    if handle_na:
        missing_cols =
df_cleaned.columns[df_cleaned.isnull().any()].tolist()
        if missing_cols:
            st.subheader("Missing Values")
            missing_df =
pd.DataFrame(df_cleaned.isnull().sum(), columns=['Missing Count'])
            st.dataframe(missing_df[missing_df['Missing
Count'] > 0])

            fill_strategy = st.sidebar.radio("Missing Value
Strategy:", ("Drop Rows", "Fill Value per Column"))
            if fill_strategy == "Drop Rows":
                if st.sidebar.button("Drop Rows with
Missing Values"):
                    df_cleaned =
cleaner.handle_missing_values(df_cleaned, strategy='drop')
                    st.session_state.data = df_cleaned
                    st.success("Rows with missing values
dropped.")

                elif fill_strategy == "Fill Value per Column":
                    fill_values = {}
                    for col in missing_cols:
                        fill_value =
st.sidebar.text_input(f"Fill '{col}' with:", value="")
                        if fill_value != "":
                            fill_values[col] = fill_value
                    if st.sidebar.button("Fill Missing
Values"):
                        if fill_values:
                            df_cleaned =
df_cleaned.fillna(fill_values)

                            st.session_state.data = df_cleaned
                            st.success("Missing values filled
per column.")

                        else:
                            st.info("Enter fill values for at
least one column.")

```

```

        else:
            st.info("No missing values found.")

        remove_duplicates = st.sidebar.checkbox("Remove
Duplicate Rows")
        if remove_duplicates:
            if st.sidebar.button("Remove Duplicates"):
                initial_rows = df_cleaned.shape[0]
                df_cleaned =
cleaner.remove_duplicate_rows(df_cleaned)
                removed_count = initial_rows -
df_cleaned.shape[0]

                st.session_state.data = df_cleaned
                st.success(f"Removed {removed_count} duplicate
rows.")

        handle_outlier = st.sidebar.checkbox("Handle Outliers")
        if handle_outlier:
            numerical_cols =
df_cleaned.select_dtypes(include=['number']).columns.tolist()
            if numerical_cols:
                column_to_clean = st.sidebar.selectbox("Column
to Handle Outliers:", numerical_cols)
                method = st.sidebar.radio("Outlier Detection
Method:", ("IQR",))

                threshold = st.sidebar.slider("IQR Threshold:",
min_value=0.5, max_value=5.0, value=1.5, step=0.1)
                outlier_handling_strategy = st.sidebar.radio(
                    "Outlier Handling Strategy:",
                    ("Mark as NaN", "Remove Outliers")
                )
                if st.sidebar.button("Handle Outliers"):
                    if method.lower() == 'iqr':
                        df_cleaned =
handle_iqr_outliers(df_cleaned, column_to_clean, threshold,
outlier_handling_strategy)

                        st.session_state.data = df_cleaned
                        st.info(f"Outliers in
'{column_to_clean}' handled.")
                    else:
                        st.info("No numerical columns available to
handle outliers.")

            if st.checkbox("Show Cleaned Data"):
                st.subheader("Cleaned Data")
                st.dataframe(st.session_state.data)

        st.sidebar.header("Ask for Insights")

```

```

        user_query = st.sidebar.text_input("Enter your question
about the data:")
        if st.sidebar.button("Get Insights"):
            if not st.session_state.data.empty:
                st.info(f"Processing your query:
'{user_query}'...")
                nl_processor = NLProcessor()
                processed_query =
nl_processor.process_query(user_query)
                st.info(f"Processed Query: {processed_query}")

                intent = processed_query['intent']
                st.info(f"Identified Intent: {intent}")

                data_selector = DataSelector()
                relevant_data =
data_selector.select_relevant_data(st.session_state.data,
processed_query['keywords'], intent)

                prompt = f"Here is the relevant
data:\n\n{relevant_data}\n\nUser question:
{processed_query['original_query']}\n\nPlease provide a concise answer
to the user's question based on this data."
                insights = generate_insights_from_llm(prompt)
                st.subheader("Generated Insights")
                st.write(insights)

                # Visualization Recommendation
                viz_recommender = VisualizationRecommender()
                recommended_visualizations =
viz_recommender.recommend_visualizations(st.session_state.data,
processed_query['keywords'], intent)

                if recommended_visualizations:
                    st.subheader("Recommended Visualizations")
                    st.write("Based on your query and the data,
here are some suggested visualizations:")
                    for viz_type in recommended_visualizations:
                        st.markdown(f"- {viz_type}")
                    else:
                        st.info("No specific visualizations
recommended for this query and data.")

                else:
                    st.warning("Please upload and (optionally)
clean the data first.")

            except Exception as e:

```

```

        st.error(f"Error loading file: {e}")

@st.cache_data
def generate_insights_from_llm(prompt: str):
    """
    Basic interaction with the Ollama DeepSeek-R1 LLM, extracting
    content outside <think> tags.
    """
    try:
        response = ollama.chat(
            model='deepseek-r1',
            messages=[
                {
                    'role': 'user',
                    'content': prompt,
                },
            ]
        )
        full_response = response['message']['content']
        # Use regular expression to find content outside <think> and
        </think>
        answer = re.sub(r'<think>.*?</think>', '', full_response,
            flags=re.DOTALL).strip()
        return answer
    except ollama.APIError as e:
        return f"Error communicating with Ollama: {e}"
    except Exception as e:
        return f"An unexpected error occurred: {e}"

def summarize_data(df: pd.DataFrame):
    """
    Generates a basic summary of the DataFrame including column names
    and data types.
    """
    summary = "Dataframe shape: {}\n\nColumns and data
types:\n".format(df.shape)
    for col in df.columns:
        summary += f"- {col}: {df[col].dtype}\n"
    return summary

def handle_iqr_outliers(df: pd.DataFrame, column_name: str, threshold:
str, strategy: str):
    """Handles outliers in a numerical column using the IQR method."""
    df_cleaned = df.copy()
    if column_name not in df_cleaned.columns or not
pd.api.types.is_numeric_dtype(df_cleaned[column_name]):
        st.error(f"Column '{column_name}' is not numeric or not
found.")
    return df_cleaned

```

```

Q1 = df_cleaned[column_name].quantile(0.25)
Q3 = df_cleaned[column_name].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - float(threshold) * IQR
upper_bound = Q3 + float(threshold) * IQR

outliers = (df_cleaned[column_name] < lower_bound) |
(df_cleaned[column_name] > upper_bound)
num_outliers = outliers.sum()
if num_outliers > 0:
    st.info(f"Identified {num_outliers} outliers in column
'{column_name}' using IQR (threshold={threshold}).")
    if strategy == "Mark as NaN":
        df_cleaned[column_name] =
df_cleaned[column_name].where(~outliers, pd.NA)
    elif strategy == "Remove Outliers":
        df_cleaned = df_cleaned[~outliers]
        st.info(f"{num_outliers} outliers removed from
'{column_name}'.")
    else:
        st.info(f"No outliers found in column '{column_name}' using IQR
(threshold={threshold}).")
        return df_cleaned

if __name__ == "__main__":
    main()

```

Insights generator

```

import pandas as pd

class DataSelector:
    def __init__(self):
        pass

    def select_relevant_data(self, df: pd.DataFrame, keywords: list,
intent: str):
        """
        Selects and returns the entire DataFrame or relevant columns as
a string for LLM analysis.
        """
        relevant_columns = [col for col in df.columns if any(keyword in
col.lower() for keyword in keywords)]

        if not relevant_columns:

```

```

        return df.to_string(index=False) # Send the entire
DataFrame as string
    else:
        return df[relevant_columns].to_string(index=False) # Send
relevant columns as string

# Example usage (can be removed later)
if __name__ == "__main__":
    data = {'product': ['A', 'B', 'A', 'C', 'B'] * 10,
            'region': ['East', 'West', 'East', 'North', 'West'] * 10,
            'sales': [100, 150, 120, 90, 160] * 10,
            'profit': [20, 30, 25, 15, 35] * 10,
            'order_date': pd.to_datetime(['2023-01-01', '2023-01-02',
'2023-01-03', '2023-01-04', '2023-01-05'] * 10)}
    df = pd.DataFrame(data)
    selector = DataSelector()
    keywords1 = ['sales', 'region']
    intent1 = "comparison"
    relevant_data1 = selector.select_relevant_data(df, keywords1,
intent1)
    print(f"Keywords: {keywords1}, Intent: {intent1}\nRelevant
Data:\n{relevant_data1[:200]}...") # Print a snippet

    keywords2 = ['profit']
    intent2 = "summary_statistics"
    relevant_data2 = selector.select_relevant_data(df, keywords2,
intent2)
    print(f"\nKeywords: {keywords2}, Intent: {intent2}\nRelevant
Data:\n{relevant_data2[:200]}...") # Print a snippet

```

## NL preprocessor

```

from nltk.corpus import stopwords # We'll still use NLTK for stopwords
(optional)

# Download stopwords if not already downloaded (run once)
try:
    stopwords.words('english')
except LookupError:
    import nltk
    nltk.download('stopwords')

class NLProcessor:
    def __init__(self):
        self.stop_words = set(stopwords.words('english')) # Using
NLTK's stopwords

        # Alternative: Define your own list of stop words if you want
to avoid NLTK entirely
        # self.stop_words = set(['the', 'a', 'an', 'is', 'are', 'in',

```

```

'on', 'at', 'which', 'what', 'how', 'why', 'of', 'for', 'to', 'and',
'but', 'or', 'if', 'then', 'else', 'when', 'where', 'who', 'whom',
'whose', 'this', 'that', 'these', 'those', 'i', 'me', 'my', 'mine',
'you', 'your', 'yours', 'he', 'him', 'his', 'she', 'her', 'hers', 'it',
'its', 'we', 'us', 'our', 'ours', 'they', 'them', 'their', 'theirs'])

def extract_keywords(self, query):
    """
    Extracts keywords from a natural language query (without NLTK's
tokenizer).
    """
    query = query.lower()
    # Replace non-alphanumeric characters (except spaces) with
spaces
    for char in "!@#$%^&*()_+~`~[]\{}|;\\':\".,/<>?":
        query = query.replace(char, ' ')
    words = query.split()
    keywords = [word for word in words if word.isalnum() and word
not in self.stop_words]
    return keywords

def identify_intent(self, keywords):
    """
    Identifies a basic intent based on keywords.
    (This is a very simple example and can be expanded)
    """
    if 'compare' in keywords or 'difference' in keywords or 'vs' in
keywords:
        return "comparison"
    elif 'trend' in keywords or 'over' in keywords and 'time' in
keywords:
        return "trend_analysis"
    elif 'summary' in keywords or 'statistics' in keywords or
'average' in keywords or 'number' in keywords:
        return "summary_statistics"
    elif 'relationship' in keywords or 'correlation' in keywords:
        return "relationship_analysis"
    elif 'tell' in keywords and 'number' in keywords: # Specific
intent example
        return "count_inquiry"
    else:
        return "general_inquiry"

def process_query(self, query):
    """
    Processes the natural language query to extract keywords and
identify intent.
    """

```

```

        keywords = self.extract_keywords(query)
        intent = self.identify_intent(keywords)
        return {"original_query": query, "keywords": keywords,
"intent": intent}

# Example usage (can be removed later):
if __name__ == "__main__":
    processor = NLProcessor()
    query1 = "Compare the sales of laptops and mice!"
    processed1 = processor.process_query(query1)
    print(f"Query: {processed1['original_query']}")
    print(f"Keywords: {processed1['keywords']}")
    print(f"Intent: {processed1['intent']}")

    query2 = "Show me the trend of revenue over time."
    processed2 = processor.process_query(query2)
    print(f"\nQuery: {processed2['original_query']}")
    print(f"Keywords: {processed2['keywords']}")
    print(f"Intent: {processed2['intent']}")

    query3 = "Tell me number of survivors?"
    processed3 = processor.process_query(query3)
    print(f"\nQuery: {processed3['original_query']}")
    print(f"Keywords: {processed3['keywords']}")
    print(f"Intent: {processed3['intent']}")

```

data cleaner

```

import pandas as pd

class DataCleaner:
    def __init__(self):
        pass

    def handle_missing_values(self, df: pd.DataFrame, strategy='drop',
fill_value=None):
        """
        Handles missing values in a Pandas DataFrame.

        Args:
            df (pd.DataFrame): The input DataFrame.
            strategy (str, optional): Strategy for handling missing
values.
                'drop': Drops rows with any missing values.
                'fill': Fills missing values with a specified value.
                Defaults to 'drop'.
            fill_value: The value to fill missing values with if
strategy is 'fill'.
                Required if strategy is 'fill'.

```



```

Returns:
    pd.DataFrame: The DataFrame with missing values handled.
    """
    df_cleaned = df.copy()
    initial_rows = df_cleaned.shape[0]
    missing_count_initial = df_cleaned.isnull().sum().sum()

    if missing_count_initial > 0:
        if strategy == 'drop':
            df_cleaned = df_cleaned.dropna()
            dropped_rows = initial_rows - df_cleaned.shape[0]
            print(f"Dropped {dropped_rows} rows with missing
values.")
        elif strategy == 'fill':
            if fill_value is not None:
                df_cleaned = df_cleaned.fillna(fill_value)
                filled_count = missing_count_initial -
df_cleaned.isnull().sum().sum()
                print(f"Filled {filled_count} missing values with:
{fill_value}")
            else:
                raise ValueError("fill_value must be provided when
strategy is 'fill'.")
        else:
            raise ValueError(f"Invalid strategy: {strategy}. Choose
'drop' or 'fill'.")
    else:
        print("No missing values found in the DataFrame.")
    return df_cleaned

def remove_duplicate_rows(self, df: pd.DataFrame):
    """
    Removes duplicate rows from a Pandas DataFrame.

    Args:
        df (pd.DataFrame): The input DataFrame.

    Returns:
        pd.DataFrame: The DataFrame with duplicate rows removed.
    """
    df_cleaned = df.copy()
    initial_rows = df_cleaned.shape[0]
    df_cleaned.drop_duplicates(inplace=True)
    removed_rows = initial_rows - df_cleaned.shape[0]
    if removed_rows > 0:
        print(f"Removed {removed_rows} duplicate rows.")
    else:
        print("No duplicate rows found.")

```

```

        return df_cleaned

# You can add more data cleaning methods here in the future
# For example:
# - handle_inconsistent_formats()
# - remove_specific_columns()
# - convert_data_types()
# - clean_text_data()

```

## Visualisation

```

import pandas as pd

class VisualizationRecommender:
    def __init__(self):
        pass

    def recommend_visualizations(self, df: pd.DataFrame, keywords:
list, intent: str):
        """
        Recommends a list of visualization types based on the data and
        intent.
        """
        recommendations = []
        numerical_cols =
df.select_dtypes(include=['number']).columns.tolist()
        categorical_cols = df.select_dtypes(include=['object',
'category']).columns.tolist()
        datetime_cols =
df.select_dtypes(include=['datetime64']).columns.tolist()

        if intent == "comparison":
            if len(numerical_cols) >= 1 and len(categorical_cols) >= 1:
                recommendations.extend(["Bar Chart", "Pie Chart"])
            elif len(numerical_cols) >= 2:
                recommendations.append("Scatter Plot")
        elif intent == "trend_analysis":
            if len(numerical_cols) >= 1 and len(datetime_cols) >= 1:
                recommendations.append("Line Chart")
        elif intent == "summary_statistics":
            if len(numerical_cols) >= 1:
                recommendations.extend(["Histogram", "Box Plot"])
                if len(numerical_cols) > 1:
                    recommendations.append("Scatter Plot (to see
distributions together)")
            elif len(categorical_cols) >= 1:
                recommendations.append("Count Plot (Bar Chart of
counts)")

```

```

        elif intent == "relationship_analysis":
            if len(numerical_cols) >= 2:
                recommendations.append("Scatter Plot")
            if len(categorical_cols) >= 1:
                recommendations.append("Grouped Scatter Plot")
        elif intent == "count_query":
            if len(categorical_cols) >= 1:
                recommendations.append("Bar Chart (Count of
categories)")

    return list(set(recommendations)) # Remove duplicates

# Example usage (can be removed later)
if __name__ == "__main__":
    data = {'product': ['A', 'B', 'A', 'C', 'B'],
            'region': ['East', 'West', 'East', 'North', 'West'],
            'sales': [100, 150, 120, 90, 160],
            'profit': [20, 30, 25, 15, 35],
            'order_date': pd.to_datetime(['2023-01-01', '2023-01-02',
'2023-01-03', '2023-01-04', '2023-01-05'])}
    df = pd.DataFrame(data)
    recommender = VisualizationRecommender()
    keywords1 = ['sales', 'region']
    intent1 = "comparison"
    recs1 = recommender.recommend_visualizations(df, keywords1,
intent1)
    print(f"Intent: {intent1}, Recommendations: {recs1}")

    keywords2 = ['trend', 'sales', 'time']
    intent2 = "trend_analysis"
    recs2 = recommender.recommend_visualizations(df, keywords2,
intent2)
    print(f"\nIntent: {intent2}, Recommendations: {recs2}")

```